

MoM: Meet - 1

20 January 2026

MEETING DEBRIEF:

- Example form types discussed included doctor-patient interaction, faculty-student conversation, registering a complaint etc. We allow the user to pick a domain and a domain-specific form to begin transcribing.
- Allow the user to create a custom domain and a custom form within it for
- The domain-specific forms have certain pre-defined set fields, we allow the user to add more fields (or remove, in general). The pipeline can suggest newer fields after the user input and processing ends.
- In case of dynamic flow of conversation as an input, allow the user to edit the form after every iteration. Allow the system to recommend edits at the end (based on performance, can be moved to mid-conversation).
- Each user shall maintain a profile consisting of all their previously-filled forms. We store all versions of the form filled by user and system as well. We could allow version-switch and multiple edits.
- Multilingual, audio and image support can be added as a layer on top of the base pipeline later into the project timeline. Initial release and data storage is confined to English.
- Input might be a running conversation among two speakers or a paragraph of monotone text. The pipeline needs to have components to distinguish among speakers, identify intent and perform contextual extraction (discard irrelevant data) alongside field-filling within a time frame of 2s.
- Due to lack of the necessity to produce benchmarks and dearth of relevant datasets, all accuracy checks will be human-review based. Sample forms, language models (Bhashini) and new field naming conventions (with other details) will be provided in later stages.
- Expected deployment is on the Product Labs servers and so will include an appropriate docker image to build and execute.
- We are not expected to make a chat interface, the chat shall be fed to the input box and our focus shall reside on optimizing the pipeline. For initial phases, we shall make our own layout and may be provided with Wireframe/UI layout details later from the client.
- Data is stored as json files with a mongodb database. Possible extension to vectordbs if the need arises.
- The client prefers the use of smaller local models like BERT, Mistral etc over API-key accessed LLMs. This point shall be kept in mind during architecture discussion - keeping a tradeoff between accuracy and latency. AI usage shall be scaled-up or down based on performance through stages.
- Conversation length is restricted to a real-life scenario of 5 minutes constituting perhaps 650-750 words.

EXPECTED DELIVERABLES:

- Dynamic chat based form filling with decent accuracy and good speed (2-3s).
- Add-ons: Voice support, Multilingual extension, Image OCR etc.
- Support for 5-10 different forms and domains.

TIMELINE:

- Overall project timeline spans from 20 January 2026 to Mid April 2026, approximately 80 days. Considerations for examinations have been requested.
- Divided into 4 phases, spanning approximately 20 days each with the first being quick with pseudo-deliverables.
- Expected meeting frequency during the course of the project - once or twice a week based on availability.

PHASE 0:

- With a timeline of 10 days, this preliminary version is expected by 1st February 2026 with 2-3 days of leeway due to impending examinations.
- Basic pipeline with support for a conversation-based form filling.
- Input shall be confined to static text and must be fed to a simple pipeline using CS scraping techniques like regex, rule-based models for field identification and form filling.
- CLI is accepted due to time constraints.
- Custom generated conversations and form templates are expected until further notice.
- Deliverables include an SRS document, Project synopsis and Architecture details. All these are subject to change with further progress as is apt.
- The architecture diagram shall use simple block diagram and flowchart notations on figma over UML as client prefers.

Further phases shall be described in detail based on upcoming meetings.