

MoM: Meet - 3

12 February 2026

MEETING DEBRIEF:

- Deployed in Product Labs Servers, that do have GPUs for significant speed-ups. The client expressed the lack of a need for speed, allowing a cap of even about 2 minutes for processing.
- Major challenges lie in ensuring smooth working and no mid-conversation crashes.

EXPECTED DELIVERABLES: (PHASE 1)

- Architecture Diagram by the end of the week, detailing a high level view of the first draft of the planned architecture.
- **Dynamic** - As the user types out the conversation, the form gets filled on the side in real-time.
- The model must recommend possible new fields that are relevant and dynamic. The user must be able to approve/deny the addition of those fields.
- Instead of typing out the entire conversation, make an “add bubble” button by setting the number of users initially. This is a UI enhancement.
- Add a “loading...” state to inform the user of the commencement of the extraction.

OTHER CHALLENGES:

- If the conversation is completely guided by a speaker, the extraction is simpler. Similarity searches may not work if the conversation is more loose.
- Having dialogues building up on previous messages must be handled, the model behavior for these cases needs to be studied and refined based on expected answers.

FUTURE DELIVERABLES:

- An account based system, allowing different users to collaborate on a conversation. This must also support real-time form filling.
- If the project ends up requiring support for voice-driven conversations, models used for ASR need to be able to determine the primary and secondary speakers. This will further drive the number/color of bubbles.